Lab 2 - Working with Structured Data

The target of this lab session is to analyze and understand a large datasets. The data we will work with today is a dataset of cities in the US and their climates. The module will discuss the challenges of loading data, finding the parts we are interested in, and visualizing data output.



The main technical tool we will be working with is a library known as Pandas. Despite the silly name, Pandas is a super popular library for data analysis. It is used in many technology companies for loading and manipulating data.

- Review

Before we get started let us review some of the Python code that we saw last class.



We first saw a bunch of different types such as numbers, strings, and lists

```
number1 = 50.5
string1 = "New York"
list1 = ["Queens", "Brooklyn", "Manhattan", "Staten Island", "The Bronx"]
```

We then saw some more complex types likes dates and counters.

```
import datetime
date1 = datetime.datetime.now()
date1

datetime.datetime(2021, 6, 28, 7, 20, 52, 750985)
```

As there are so many different types in Python, we discussed how important it was to use Google and StackOverflow to find examples.

```
from collections import Counter
counter = Counter(["A", "B", "A", "A", "C", "C", "B", "A", "A", "A"])
counter.most_common()
    [('A', 6), ('B', 2), ('C', 2)]
```

Next, we focused on if and for the two most important control elements in python.

- if lets us decide which block of code to run
- for lets us run the same code for each element in a list

```
for val in range(10):
    print(val)

    0
    1
    2
    3
    4
    5
    6
    7
    8
    9
```

Finally we discussed the special case of strings. There are many useful ways to find values in strings and create new strings.

```
"first " + "second"
    'first second'
str1 = "first|second"
```

```
str1.split("|")
    ['first', 'second']
```

Questions from the previous lab?

Review Exercise

Print only the values in this list that are greater than 20.

→ Unit A

This week is all about datatables, also known as dataframes. Data tables are a common way of representing facts anywhere from newspaper articles to scientific studies.

For instance, as a running example let us consider this table from Wikipedia.



https://en.wikipedia.org/wiki/List_of_North_American_cities_by_population

You may have used datatables before in spreadsheets. For example we can put that wikipedia table in a spreadsheet.

https://docs.google.com/spreadsheets/d/1Jwcr6IBJbOT1G4Vq7VqaZ7S1V9gRmUb5ALkJPaG5fxI/edit?usp=sharing

In this spreadsheet we can do lots of things.

Student question: Do you know how to do the following?

- · Change the column names
- · Delete a row
- · Make a graph
- Add a new column

What about more advanced ideas. Can you?

- · Sort by a column?
- Add a new column that changes the previous column?
- Take the sum of a row?
- Find the highest value in a row?

In this lab we will work with real-world data to learn how to calculate important properties.

▼ Pandas

The data that we are working with is located in the file "Cities.csv". This file is raw data as a text file. We can see the output in raw form.

https://srush.github.io/BT-AI/notebooks/Cities.csv

We can see that "csv" stands for "comma separated values" as each element of the file is split using a comma.

Pandas is as a super-powered spreadsheet.

import pandas as pd

To load data in the library we use the following command. Here df refers to the "DataFrame" which is what Pandas calls a spreadsheet.

df = pd.read_csv("https://srush.github.io/BT-AI/notebooks/Cities.csv")
df

	Rank	City	Country	Population
0	0	Mexico City	Mexico	8918653
1	1	New York City	United States	8550405
2	2	Los Angeles	United States	3971883
3	3	Toronto	Canada	2826498
4	4	Chicago	United States	2720546
90	90	Surrey	Canada	526004
91	91	Ciudad López Mateos	Mexico	523296
92	92	Tultitlán	Mexico	520557
93	93	Fresno	United States	520052
94	94	Carrefour	Haiti	501768

95 rows × 4 columns

Just like in a spreadsheet Pandas has multiple columns representing the underlying elements in the data. These each have a name here.

```
df.columns
```

```
Index(['Rank', 'City', 'Country', 'Population'], dtype='object')
```

To see just the elements in a single column we can use square brackets to see just only column.

```
df["City"]
```

```
0
              Mexico City
1
            New York City
2
               Los Angeles
3
                   Toronto
                   Chicago
90
                    Surrey
      Ciudad López Mateos
91
92
                 Tultitlán
93
                    Fresno
                 Carrefour
```

Name: City, Length: 95, dtype: object

Student Question: Can you print out another column in the table?

```
# 🕅 🕅 🕅 FILLME
df["Country"]
     0
                  Mexico
     1
           United States
     2
           United States
     3
                  Canada
     4
           United States
     90
                  Canada
     91
                  Mexico
     92
                  Mexico
     93
           United States
     94
                   Haiti
     Name: Country, Length: 95, dtype: object
```

Alternatively if we want to see just a single row we can use the loc command.

df.loc[1]

Rank 1
City New York City
Country United States
Population 8550405
Name: 1, dtype: object

If we want to select several rows we can also pass in a list.

	Rank	City	Country	Population
1	1	New York City	United States	8550405
5	5 5	Houston	United States	2296224
6	6	Havana	Cuba	2117625

Student Question: Can you print out the rows of Philadelphia and Los Angeles?

#2 2 2 FILLME df.loc[[5,9]]

	Rank	City	Country	Population
5	5	Houston	United States	2296224
9	9	Philadelphia	United States	1567442

▼ Filters

These commands are relatively basic though and easy to do in a standard spreadsheet. The main power of Pandas comes from the ability to select rows based on complex filters.

For instance, if you were in a spreadsheet, how would you select only the rows that correspond to cities in *Mexico*? It's possible but a bit challenging.

In Pandas, we first create a filter. This is kind of like an if statement that gets applied to every row. It creates a variable that remembers which rows passed the filter test.

Filtering

- 1. Decide on the conditional statements in your filter.
- 2. Define a filter varaible for your dataframe.
- 3. Apply filter and rename the dataframe.

Step 1. Our filter is that we want the Country column to be Mexico

Step 2. We create a filter variable with this conditional. Notice that the filter has a 1 for every city in Mexico and a 0 otherwise.

```
filter = df["Country"] == "Mexico"
filter
     0
            True
     1
           False
     2
           False
     3
           False
           False
     90
           False
     91
            True
     92
            True
     93
           False
     94
           False
     Name: Country, Length: 95, dtype: bool
```

Step 3. We then apply the filter to select the rows that we would like to keep around.

cities_in_mexico_df = df.loc[filter]
cities_in_mexico_df

	Rank	City	Country	Population
0	0	Mexico City	Mexico	8918653
8	8	Ecatepec de Morelos	Mexico	1677678
12	12	Guadalajara	Mexico	1460148
13	13	Puebla	Mexico	1437939
15	15	Juárez	Mexico	1382753
16	16	León	Mexico	1349224
18	18	Tijuana	Mexico	1298475
21	21	Zapopan	Mexico	1179681
22	22	Monterrey	Mexico	1109171
24	24	Nezahualcóyotl	Mexico	1039867
29	29	Naucalpan	Mexico	970012

We need to be careful to give this a new name. It does not change the original dataframe it just shows us the rows we asked for.

```
AT 05 TI 14 ' 070500
```

Filtering is a really important step because it lets us calculate other properties.

For example, we can then count the number of cities in Mexico.

```
total_cities_in_mexico = cities_in_mexico_df["City"].count()
total_cities_in_mexico
```

Or we can count the population of the biggest cities in Mexico.

Filters can also be more complex. You can check for any of the different properties you might check for in a standard if statement.

For instance, here we want to keep both cities in the US and in Canada. The symbol | means either-or.

https://colab.research.google.com/drive/112 X-XbTy6mHFneWTPYFhnakf2l08bq#scrollTo=6da1d0dc&printMode=true

78 78 Morelia Mexico 608190

filter = (df["Country"] == "United States") | (df["Country"] == "Canada")
us_or_canada_df = df.loc[filter]
us_or_canada_df

	Rank	City	Country	Population
1	1	New York City	United States	8550405
2	2	Los Angeles	United States	3971883
3	3	Toronto	Canada	2826498
4	4	Chicago	United States	2720546
5	5	Houston	United States	2296224
7	7	Montreal	Canada	1753034
9	9	Philadelphia	United States	1567442
10	10	Phoenix	United States	1563025
11	11	San Antonio	United States	1469845
14	14	San Diego	United States	1394928
17	17	Dallas	United States	1300092
19	19	Calgary	Canada	1230915
25	25	San Jose	United States	1026908
30	30	Ottawa	Canada	956710
31	31	Austin	United States	931830
32	32	Edmonton	Canada	899447
36	36	Jacksonville	United States	868031
38	38	San Francisco	United States	864816

Student Question: How many of the cities are in the US or Canada?

40 01 1 11 11 101 1 00000

46

46 46 Mississauga Canada /61300

Here is a list of the different statements that we commonly use.

57	57	Seattle	I Inited 9	States Filter	684451 Symbol
			(Or	1
			,	And	&
			1	Not	~
			1	Equal	==
			1	Less	<
			(Greater	>

	Filter	Symbol
	Greater	>
	In	.str.contains
United		.isin
	United	Greater

Note: I didn't know many of these by heart. Don't be afraid to google "how to filter by ... in pandas" if you get stuck.

→ Group Exercise A

▼ Question 1

71

05 05 11---:14--- 0----1- 550050

Filters can be of many different types. For instance, when working with numerical fields we can have filters based on greater-than and less-than comparisons.

Write a filter that keeps only cities with greater than a million people.

```
93 Fresno United States 520052
```

	Rank	City	Country	Population
0	0	Mexico City	Mexico	8918653
1	1	New York City	United States	8550405
2	2	Los Angeles	United States	3971883
3	3	Toronto	Canada	2826498
4	4	Chicago	United States	2720546
5	5	Houston	United States	2296224
6	6	Havana	Cuba	2117625
7	7	Montreal	Canada	1753034
8	8	Ecatepec de Morelos	Mexico	1677678
9	9	Philadelphia	United States	1567442
10	10	Phoenix	United States	1563025
11	11	San Antonio	United States	1469845
12	12	Guadalajara	Mexico	1460148
13	13	Puebla	Mexico	1437939
4 4	11	San Diago	United States	120/1020

How many are there?

```
# 📝 📝 📝 📝 FILLME
```

row_index = df["Population"].argmin()
df.loc[row_index]

Rank 94
City Carrefour
Country Haiti
Population 501768
Name: 94, dtype: object

Mantagray Marian 440047

(Be sure to print it out to check that it worked!)

▼ Question 2

26 26 Santo Domingo Dominican Republic 1007997

Several cities in North America include the word "City" in their name. Write a filter to find the cities that have "City" in their name.

```
tilter = dt["City"].str.contains("City")
city_name = df.loc[filter]
city_name
```

	Rank	City	Country	Population
0	0	Mexico City	Mexico	8918653
1	1	New York City	United States	8550405
27	27	Guatemala City	Guatemala	994078
73	73	Oklahoma City	United States	631346
86	86	Quebec City	Canada	540994

What is the smallest city on this list?

```
# | FILLME df.nsmallest
```

≺bou	nd me	thod DataFrame.nsmall	est of	Rank		City	Country	Popu]
0	0	Mexico City		Mexico	8918653			
1	1	New York City	United	States	8550405			
2	2	Los Angeles	United	States	3971883			
3	3	Toronto		Canada	2826498			
4	4	Chicago	United	States	2720546			
		• • •			• • •			
90	90	Surrey		Canada	526004			
91	91	Ciudad López Mateos		Mexico	523296			
92	92	Tultitlán		Mexico	520557			
93	93	Fresno	United	States	520052			
94	94	Carrefour		Haiti	501768			
[95	rows	x 4 columns]>						
[]]	. 0115	X . columns],						
4								•

▼ Question 3

Most of the cities on the list are in Canada, Mexico or the US.

Can you write a filter to find the cities that are not in any of these countries?

```
# D D D FILLME

filter = df["City"].isnull()

city_name = df.loc[filter]

city_name
```

Rank City Country Population

What is the largest city in this list?

▼ Question 4

We can also apply filters that look for two properties at the same time.

Can you write a filter to find the cities in the US of over a million people?

```
# | FILLME
filter2 = df["Population"] > 1,000,000
filter2
            True
     (0
      1
            True
            True
      3
            True
            True
            . . .
      90
            True
      91
            True
      92
            True
      93
            True
      94
            True
      Name: Population, Length: 95, dtype: bool, 0, 0)
```

How many are there?

```
# | FILLME

populations_num_df = df.loc[filter3]

print(populations_num_df.iloc[-1])
```

→ Unit B

In this unit we will look at three more advanced Pandas functions. Unlike filters, which just remove rows, these will allow use to manipute our data to compute new properties and even new columns.

▼ Group By's

We saw above how to compute the total number of cities in Mexico on our list. We did this by first filtering and then "aggregating" by calling <code>count()</code>. Here is a reminder.

```
filter = df["Country"] == "Mexico"
cities_in_mexico_df = df.loc[filter]
total_cities_in_mexico = cities_in_mexico_df["City"].count()
total_cities_in_mexico
```

However, what if we also want to know the number of cities in Canada and US and all the other countries on our list. We can do this with a group-by operation

GroupBy

- 1. GroupBy Determine the subset of data to use
- 2. Aggregation Compute a property over the group

```
Step 1. Group By
```

```
grouped = df.groupby(["Country"], as_index=False)
```

Step 2. Aggregate

count_of_cities = grouped["City"].count()
count_of_cities

	Country	City
0	Canada	12
1	Cuba	1
2	Dominican Republic	2
3	Guatemala	2
4	Haiti	2
5	Honduras	2
6	Jamaica	1
7	Mexico	38
8	Nicaragua	1
9	United States	34

Here is another example. This one computes the population of the largest city in each country.

	Country	Population
0	Canada	2826498
1	Cuba	2117625
2	Dominican Republic	1007997
3	Guatemala	994078
4	Haiti	987310
5	Honduras	1190230
6	Jamaica	669627
7	Mexico	8918653
8	Nicaragua	1048134
9	United States	8550405

** Student Question: Can you compute the city with the minimum population on the list for each country? **

```
# | FILLME
min_pop = grouped["Population"].min()
min_pop
```

	Country	Population
0	Canada	526004
1	Cuba	2117625
2	Dominican Republic	723162
3	Guatemala	564686
4	Haiti	501768
5	Honduras	742118
6	Jamaica	669627
7	Mexico	520557
8	Nicaragua	1048134
9	United States	520052

Manipulating Tables

Another useful aspect of tables is is to add in new columns. Adding new columns allows us to group by additional properties, create advanced filters, or make pretty graphs.

The easiest way to add a new column in pandas is to write a function that tells us how to create the new column from the other columns in the table.

In order to add a new column, we need to write a function. If you remember last class, a function looked something like this.

```
# Returns if the country is in US or Canada
def in_us_or_canada(country):
    if country == "United States":
        return "US/Canada"
    if country == "Canada":
        return "US/Canada"
```

```
return "Not US/Canada"
print(in_us_or_canada("Mexico"))
print(in_us_or_canada("Canada"))
Not US/Canada
```

US/Canada

Now we can add a new column by setting that column equal to the country. We do this by calling Pandas map with the function and the column of interest. This line of code will call our function for each row of the Country column. Notice how it creates a new column.

	Rank	City	Country	Population	US_or_Canada
0	0	Mexico City	Mexico	8918653	Not US/Canada
1	1	New York City	United States	8550405	US/Canada
2	2	Los Angeles	United States	3971883	US/Canada
3	3	Toronto	Canada	2826498	US/Canada
4	4	Chicago	United States	2720546	US/Canada
90	90	Surrey	Canada	526004	US/Canada
91	91	Ciudad López Mateos	Mexico	523296	Not US/Canada
92	92	Tultitlán	Mexico	520557	Not US/Canada
93	93	Fresno	United States	520052	US/Canada
94	94	Carrefour	Haiti	501768	Not US/Canada

95 rows × 5 columns

```
df.columns
```

```
Index(['Rank', 'City', 'Country', 'Population', 'US_or_Canada'], dtype='object')
```

We can then use this column in a group by.

```
grouped = df.groupby(["US_or_Canada"])
count_of_cities = grouped["City"].count()
count_of_cities
```

```
US_or_Canada
Not US/Canada 49
US/Canada 46
Name: City, dtype: int64
```

A similar technique can be used to manipulate the data in a column to change certain values. For instance, we might want to remove the final "City" from cities like "New York"

```
def change_name(str1):
    return str1.replace(" City", "")
change_name("New York City")
    'New York'

df["City"] = df["City"].map(change_name)
df
```

Rank		City	Country	Population	US_or_Canada
0	0	Mexico	Mexico	8918653	Not US/Canada
1	1	New York	United States	8550405	US/Canada
2	2	Los Angeles	United States	3971883	US/Canada
3	3	Toronto	Canada	2826498	US/Canada
4	4	Chicago	United States	2720546	US/Canada
90	90	Surrey	Canada	526004	US/Canada
91	91	Ciudad López Mateos	Mexico	523296	Not US/Canada
92	92	Tultitlán	Mexico	520557	Not US/Canada
93	93	Fresno	United States	520052	US/Canada
94	94	Carrefour	Haiti	501768	Not US/Canada

95 rows × 5 columns

▼ Joining Together Tables

Pandas becomes much more powerful when we start to have many different tables that relate to each other. For this example we will consider another table that provides the locations about these cities. You can see that here:

City Location Spreadsheet

Lets load this table into a new variable.

	Id	City	Country	Longitude	Latitude
0	0	A Coruña	Spain	8.73W	42.59N
1	1	Aachen	Germany	6.34E	50.63N
2	2	Aalborg	Aalborg Denmark		57.05N
3	3	Aba	Nigeria	8.07E	5.63N
4	4	Abadan	Iran	48.00E	29.74N
3505	3505	Århus	Denmark	10.33E	57.05N
3506	3506	Çorlu	Turkey	27.69E	40.99N
3507	3507	Çorum	Turkey	34.08E	40.99N
3508	3508	Öskemen	Kazakhstan	82.39E	50.63N
3509	3509	Ürümqi	China	87.20E	44.20N

3510 rows × 5 columns

This table has most of the cities in our dataset. But there are also a lot of other cities in this table outside of North America.

```
filter = all_cities_df["Country"] == "Germany"
europe_df = all_cities_df.loc[filter]
europe df
```

	Id	City	Country	Longitude	Latitude
1	1	Aachen	Germany	6.34E	50.63N
187	187	Augsburg	Germany	10.66E	47.42N
338	338	Bergisch Gladbach	Germany	6.34E	50.63N
340	340	Berlin	Germany	13.14E	52.24N
370	370	Bielefeld	Germany	7.88E	52.24N

In order to use this new information let's merge since it in to our table. We just need to tell pandas which are the shared columns between the two tables.

```
df = df.merge(all_cities_df, on=["City", "Country"])
df
```

	Rank	City	Country	Population	US_or_Canada	Id	Longitude	Latitude
0	0	Mexico	Mexico	8918653	Not US/Canada	1955	98.96W	20.09N
1	1	New York	United States	8550405	US/Canada	2126	74.56W	40.99N
2	2	Los Angeles	United States	3971883	US/Canada	1775	118.70W	34.56N
3	3	Toronto	Canada	2826498	US/Canada	3140	80.50W	44.20N
4	4	Chicago	United States	2720546	US/Canada	608	87.27W	42.59N
79	87	Tonalá	Mexico	536111	Not US/Canada	3132	104.08W	20.09N
80	88	Tucson	United States	531641	US/Canada	3171	111.20W	31.35N
81	89	Cuautitlán · ···	Mexico	531041	Not	719	98.96W	20.09N

→ Group Exercise B

▼ Question 1

The following website has the official abbreviation codes for all countries. https://en.wikipedia.org/wiki/List_of_ISO_3166_country_codes To get all of the unique entries of a series you can use https://pandas.pydata.org/docs/reference/api/pandas.unique.html

Using a table mapping countries to their codes. Can you add a new column to the table called "Abbrev" that lists the abbreviation code for that city?

```
def abbreviate(country):
   if country == "United States":
       return "US"
   if country == "Canada":
       return "CA"
   if country == "Mexico":
       return "MX"
   if country == "Cuba":
       return "CU"
   if country == "Dominican Republic":
       return "DR"
   if country == "Guatemala":
       return "GT"
   if country == "Haiti":
       return "HT"
   if country == "Honduras":
       return "HN"
   if country == "Jamaica":
       return "JM"
   if country == "Nicaragua":
       return "NI"
```

▼ Question 2

Our table has the Latitude and Longitude of all the major North American Cities.

Can you find out where New York is located? How about Detroit, Las Vegas, and Portland?

▼ Question 3

Currently in the table the latitude and longitude are represented as string types, because they have

Instead we would for longitude E to be positive and W negative, for example 100W -> -100 For latitude we would like N to be positive and S to be negative, for example 100N -> 100

```
def latitude to number(latitude string):
   if "N" in latitude string:
      latitude string = latitude string.replace("N", "")
      return int(latitude string)
   if "S" in latitude string :
      latitude string = latitude string.replace("S", "")
      return 0-int(latitude string)
def longitude_to_number(longitude_string):
   if "E" in latitude string:
      latitude string = latitude string.replace("E", "")
      return int(latitude string)
   if "W" in latitude string :
      latitude_string = latitude_string.replace("W", "")
      return 0-int(latitude string)
lat = latitude to number("190N")
lat
     190
```

Can you use these functions to fix the Latitude and Longitude columns to instead use numeric values?

```
# | FILLME

df.groupby(['Latitude', 'Longitude']).sum()
```

		капк	Population	10
Latitude	Longitude			
12.05N	85.48W	23	1048134	1862
13.66N	86.70W	20	1190230	3062

▼ Question 4

40.001 400.4014 55 740004 40

After completing question 3 use group by and compute the Latitude of most southern city in each country of the table.

TITLET 1818TT 00 010001 2010

#②②② FILLME df.groupby(['Latitude']).min()

		Rank	City	Country	Population	US_or_Canada	Id	Longitude
La	atitude							
•	12.05N	23	Managua	Nicaragua	1048134	Not US/Canada	1862	85.48W
1	13.66N	20	Tegucigalpa	Honduras	1190230	Not US/Canada	3062	86.70W
1	15.27N	27	Guatemala	Guatemala	564686	Not US/Canada	1079	87.50W
						K I = 1		

Visualization

Next class we will dive deeper into plotting and visualization. But let's finish with a little demo to show off all the tables we created.

```
21.70N
                  16 Aguascalientes
                                             Mexico
                                                         /11682
                                                                                        101.25W
                                                                                  39
                                                                    IIC/Canada
First we import some libraries
       _ ..._..
                                                                    US/Canada
import altair as alt
from vega datasets import data
                                                                           K I _ £
states = alt.topo feature(data.us 10m.url, feature='states')
background = alt.Chart(states).mark_geoshape().project('albersUsa')
       43.14IN
                              Ausun
                                             IVIEXICO
                                                         013333
                                                                                 ıσı
                                                                                        110.// 00
                                                                    US/Canada
Now we can plot
states = alt.topo feature(data.world 110m.url, feature='countries')
chart = alt.Chart(states).mark_geoshape(
        fill='lightgray',
        stroke='white'
    ).properties(
        width=500,
        height=300
    ).project('orthographic', rotate= [95, -42, 0])
if True:
    points = alt.Chart(df).mark_circle().encode(
        longitude='Longitude',
        latitude='Latitude',
```

size="Population",

chart += points

chart

tooltip=['City', 'Population']

```
SchemaValidationError
                                          Traceback (most recent call last)
/usr/local/lib/python3.7/dist-packages/altair/vegalite/v4/api.py in to_dict(self,
*args, **kwargs)
    380
                if dct is None:
                    kwargs["validate"] = "deep"
    381
--> 382
                    dct = super(TopLevelMixin, copy).to_dict(*args, **kwargs)
    383
                # TODO: following entries are added after validation. Should they be
    384
validated?
                                   16 frames -
/usr/local/lib/python3.7/dist-packages/altair/utils/schemapi.py in to_dict(self,
validate, ignore, context)
    337
                        self.validate(result)
    338
                    except jsonschema.ValidationError as err:
--> 339
                        raise SchemaValidationError(self, err)
    340
                return result
    341
SchemaValidationError: Invalid specification
        altair.vegalite.v4.schema.channels.Latitude->type, validating 'enum'
        'nominal' is not one of ['quantitative']
```



×

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.